

[Skip to main content](#)

- [Home](#)
- [Topics](#)
- [Reference](#)
- [Glossary](#)
- [Help](#)
- [Notebook](#)

Virtual Workshop

Welcome guest

[Log in \(Globus\)](#)

[Log in \(other\)](#)

[Try the quiz before you start](#)

MPI Collective Communications

[Introduction](#) [Goals](#) [Prerequisites](#)

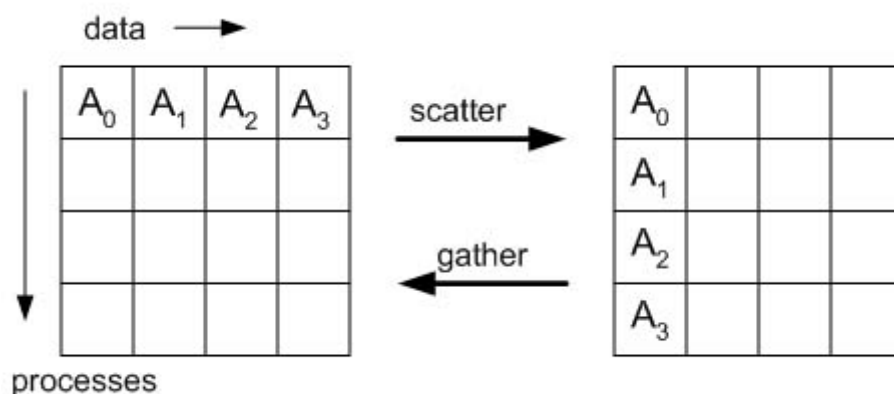
[Characteristics](#) [Three Types of Routines](#) [Barrier Synchronization](#) [Data Movement](#) • [Broadcast](#) • [Gather and Scatter](#) • [Gather/Scatter Effect](#) • [Gatherv and Scatterv](#) • [Allgather](#) • [All to All](#) [Global Computing](#) • [Reduce](#) • [Scan](#) • [Operations and Example](#) • [Allreduce Mini-Exercise](#) [Nonblocking Routines](#) • [Nonblocking Example](#) [Performance Issues](#) • [Two Ways to Broadcast](#) • [Two Ways to Scatter](#) [Application Example](#) • [Scatter vs. Scatterv](#) • [Scatterv Syntax](#)

[Exercise Quiz](#)

[Short survey](#)

MPI Collective Communications: Gather/Scatter Effect

In order to illustrate the gather and scatter functions, we give a matrix-style depiction below:



Let's consider in detail how MPI_Gather might be used to facilitate a distributed matrix computation.

Example: matrix-vector multiplication

- Matrix is distributed by rows (i.e. row-major order)
- Product vector is needed in entirety by one process
- MPI_Gather will be used to collect the product from processes

Description of Sample Code

The problem associated with the following sample code is the multiplication of a matrix A, size 100x100, by a vector b of length 100. The example uses four MPI processes, so each process will work on its own chunk of 25 rows of A. Since b is the same for each process, it will simply be replicated across processes. The vector c will therefore have 25 elements calculated by each process; these are stored in cpart. Here is a picture of how the overall computation is distributed:

$$\begin{array}{c}
 \text{A} \quad * \quad \text{b} \quad = \quad \text{c} \\
 \left[\begin{array}{c} \text{Process 0} \\ \text{-----} \\ \text{Process 1} \\ \text{-----} \\ \text{Process 2} \\ \text{-----} \\ \text{Process 3} \end{array} \right] \left[\begin{array}{c} \\ \\ \\ \end{array} \right] = \left[\begin{array}{c} 0 \\ \text{---} \\ 1 \\ \text{---} \\ 2 \\ \text{---} \\ 3 \end{array} \right]
 \end{array}$$

A: matrix distributed by rows

b: vector shared by all processes

c: vector updated by each process independently

The MPI_Gather routine will retrieve cpart from each process and store the result in ctotol, which is the complete vector c.

Sample Code in C

```
double a[25,100],b[100],cpart[25],ctotal[100];
int root;
root=0;
for(i=0;i<25;i++)
{
    cpart[i]=0;
    for(k=0;k<100;k++)
    {
        cpart[i]=cpart[i]+a[i,k]*b[k];
    }
}
MPI_Gather(cpart,25,MPI_REAL,ctotal,25,MPI_REAL,root,
          MPI_COMM_WORLD);
```

Sample Code in FORTRAN

```
! Fortran, unlike C, stores matrices in column-major order,
! so the A below is the transpose of the actual A matrix.
DIMENSION A(100,25), b(100), cpart(25), ctotol(100)
INTEGER root
DATA root/0/
DO I=1,25
    cpart(I) = 0.
    DO K=1,100
```

```
        cpart(I) = cpart(I) + A(K,I)*b(K)
    END DO
END DO
CALL MPI_GATHER(cpart, 25, MPI_REAL, ctotal, 25, MPI_REAL, &
    root, MPI_COMM_WORLD, ierr)
```

[<= previous](#)[next =>](#)

© 2021 [Cornell University](#) | [Cornell University Center for Advanced Computing](#) | [Copyright Statement](#) | [Terminology Statement](#)