

[Skip to main content](#)

---

- [Home](#)
- [Topics](#)
- [Reference](#)
- [Glossary](#)
- [Help](#)
- [Notebook](#)

# Virtual Workshop

Welcome guest

[Log in \(Globus\)](#)

[Log in \(other\)](#)

[Try the quiz before you start](#)

MPI Collective Communications

[Introduction](#) [Goals](#) [Prerequisites](#)

[Characteristics](#) [Three Types of Routines](#) [Barrier Synchronization](#) [Data Movement](#) • [Broadcast](#) • [Gather and Scatter](#) • [Gather/Scatter Effect](#) • [Gatherv and Scatterv](#) • [Allgather](#) • [All to All](#) [Global Computing](#) • [Reduce](#) • [Scan](#) • [Operations and Example](#) • [Allreduce Mini-Exercise](#) [Nonblocking Routines](#) • [Nonblocking Example](#) [Performance Issues](#) • [Two Ways to Broadcast](#) • [Two Ways to Scatter](#) [Application Example](#) • [Scatter vs. Scatterv](#) • [Scatterv Syntax](#)  
[Exercise Quiz](#)  
[Short survey](#)

---

## MPI Collective Communications: Gather and Scatterv

MPI\_Gatherv and MPI\_Scatterv are the variable-message-size versions of MPI\_Gather and MPI\_Scatter. MPI\_Gatherv extends the functionality of MPI\_Gather to permit a varying count of data from each process, and to allow some flexibility in where the gathered data is placed on the root process. It does this by changing the count argument from a single integer to an integer array and providing a new argument `displs` (an array). MPI\_Scatterv extends MPI\_Scatter in a similar manner. More information on the use of these routines will be presented in an [Application Example](#) later in this module.

## C

```
int MPI_Gatherv(const void* sbuf, int scount, \
               MPI_Datatype stype, void* rbuf, const int rcounts[], \
               const int displs[], MPI_Datatype rtype, \
               int root, MPI_Comm comm)

int MPI_Scatterv(const void* sbuf, const int scounts[], \
                const int displs[], MPI_Datatype stype, void* rbuf, \
                int rcount, MPI_Datatype rtype, \
                int root, MPI_Comm comm)
```

## FORTRAN

```
MPI_GATHERV(sbuf, scount, stype, rbuf, rcounts, displs, rtype,
            root, comm, ierr)
```

```
MPI_SCATTERV(sbuf, scounts, displs, stype, rbuf, rcount, rtype,  
             root, comm, ierr)
```

The variables for **Gatherv** are:

- sbuf** starting address of send buffer,
- scount** number of elements in send buffer,
- stype** data type of send buffer elements,
- rbuf** starting address of receive buffer,
- rcounts** array containing number of elements to be received from each process,
- displs** array specifying the displacement relative to rbuf at which to place the incoming data from corresponding process,
- rtype** data type of receive buffer,
- root** rank of receiving process,
- comm** group communicator.

Note: rbuf, rcounts, displs, rtype are significant for the root process only.


The variables for **Scatterv** are:

- sbuf** address of send buffer,
- scounts** integer array specifying the number of elements to send to each process,
- displs** array specifying the displacement relative to sbuf from which to take the data going out to the corresponding process,
- stype** data type of send buffer elements,
- rbuf** address of receive buffer,
- rcount** number of elements in receive buffer,
- rtype** data type of receive buffer elements,
- root** rank of sending process,
- comm** group communicator

Note: sbuf, scounts, displs, stype are significant for the root process only.

For the purpose of illustrating the usage of MPI\_Gatherv and MPI\_Scatterv, we give two Fortran program fragments below:

## MPI\_GATHERV Example




```

real a(25), rbuf(MAX)
integer displs(NX), rcounts(NX), nsize
...
do i= 1, nsize
    displs(i) = (i-1) * stride
    rcounts(i) = 25
enddo
call mpi_gatherv(a, 25, MPI_REAL, rbuf, rcounts, displs, &
    MPI_REAL, root, comm, ierr)
...

```

Notice that the effect of setting stride= 37.5 is to cause the separation between the chunks on the root to alternate between 12 and 13 integers.

## MPI\_SCATTERV Example



```

real a(25), sbuf(MAX)
integer displs(NX), scounts(NX), nsize
...
do i= 1, nsize
    displs(i) = (i-1) * stride
    scounts(i) = 25
enddo
call mpi_scatterv(sbuf, scounts, displs, MPI_REAL, a, 25, &
    MPI_REAL, root, comm, ierr)
...

```

Notice again that the effect of setting  $\text{stride}=37.5$  is to cause the separation between the chunks on the root to alternate between 12 and 13 integers.

[<= previous](#)[next =>](#)

---

© 2021 [Cornell University](#) | [Cornell University Center for Advanced Computing](#) | [Copyright Statement](#) | [Terminology Statement](#)